

Deep Learning Topic Based Sentiment Analysis

Human Language Technologies
2019 / 2020

Berti Stefano

Abstract

The aim of this project is to apply the model

<https://github.com/cbaziotis/datastories-semeval2017-task4>

to the Aspect Category Polarity task of the Absita competition

<http://sag.art.uniroma2.it/absita>

To show the strength of the proposed model. I will try two different word vectors strategies: the Italian Word2Vec

<https://mlunicampania.gitlab.io/italian-word2vec/>

and the Italian BERT AlBERTo

<http://ceur-ws.org/Vol-2481/paper57.pdf>

and I will also try to apply the same model to the Aspect Category Detection task.

1 Task, dataset and metrics

The Absita competition is divided into 2 tasks:

- **ACD**: Aspect Category Detection, to understand which topic is dealt inside the review
- **ACP**: Aspect Category Polarity, given a review and a topic, to understand if the topic is dealt in a positive, negative, neutral or mixed way

The second task is dependent from the first one, but we will see it as a independent task because, otherwise, the output of the first task would influence the input of the first task, and so the second task results couldn't be better than the results of the first task. I transformed the given dataset in order to obtain a tsv file of the form

id, topic, y, review

where

- **id** is the id of the review
- **topic** is one element in ['cleanliness', 'amenities', 'value', 'wifi', 'location', 'staff', 'other']
- **y** in ACP task, it refers to the sentiment of the review towards that topic and it is an element in ['positive', 'negative', 'neutral', 'mixed'], in ACD

task it refers to if the topic is dealt in the review or not and it is an element in ['positive', 'negative']

- **review** is the raw review

Table 1: number of element for class

data	positive	neutral	mixed	negative
train	4942	0	173	3797
test	2080	0	64	1757

I repeated the reviews with more than one topic in order to have one entry for each topic. Since I didn't have any element for neutral review neither in the train set nor in the test set, I decided to remove it from the possible classes, because I didn't have any way to explain to my model how a neutral review is like. The metrics used in this competition are micro-precision, micro-recall and f1-score.

2 Description of the models

The model chosen is the one who outperformed *Semeval2017*, which tasks were very similar to this competition. It is a deep-learning model with context-aware attention:

- The inputs are word indices of topics and reviews to be embedded in the EmbeddingLayer for the Word2Vec version or topics and reviews already embedded for the ALBERTo version
- It feeds the reviews' embedding in a shared LSTM. It does the same with the topic's embedding using the same weights in order to try to get meaningful representations.
- It concatenates each word representation with the topic representation.
- It uses a context-aware attention mechanism, which tries to understand which part of the reviews contribute more to understand better sentiment/references towards the topic
- It feeds those representations in a dense layer with a single sigmoid neuron for task ACD and 3 softmax neurons for task ACP

2.1 ACP

Initially I only considered the positive and negative classes because of the very few mixed samples, obtaining from mixed reviews a random sentiment, and, although this gives nice results, I moved from one sigmoid output neuron for

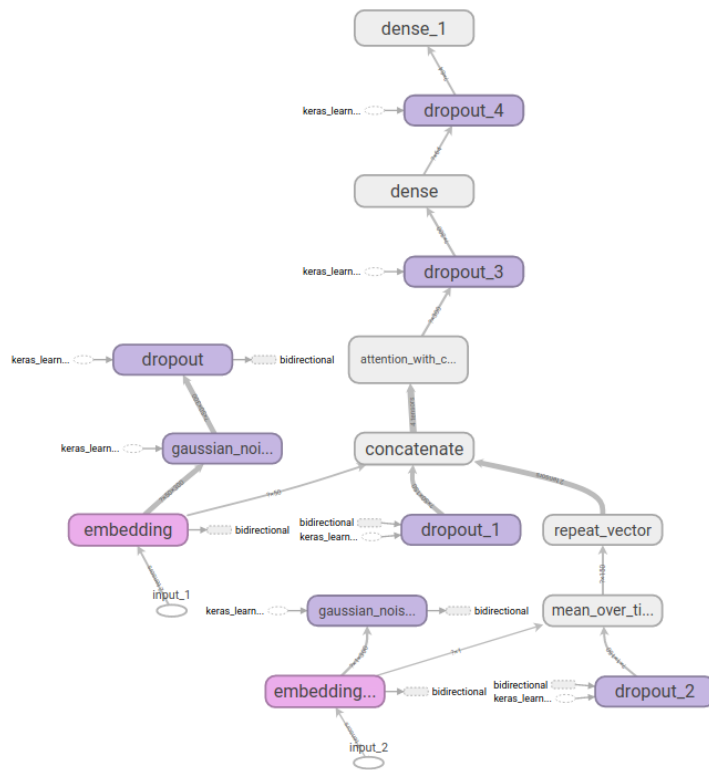


Figure 1: The model that uses Word2Vec

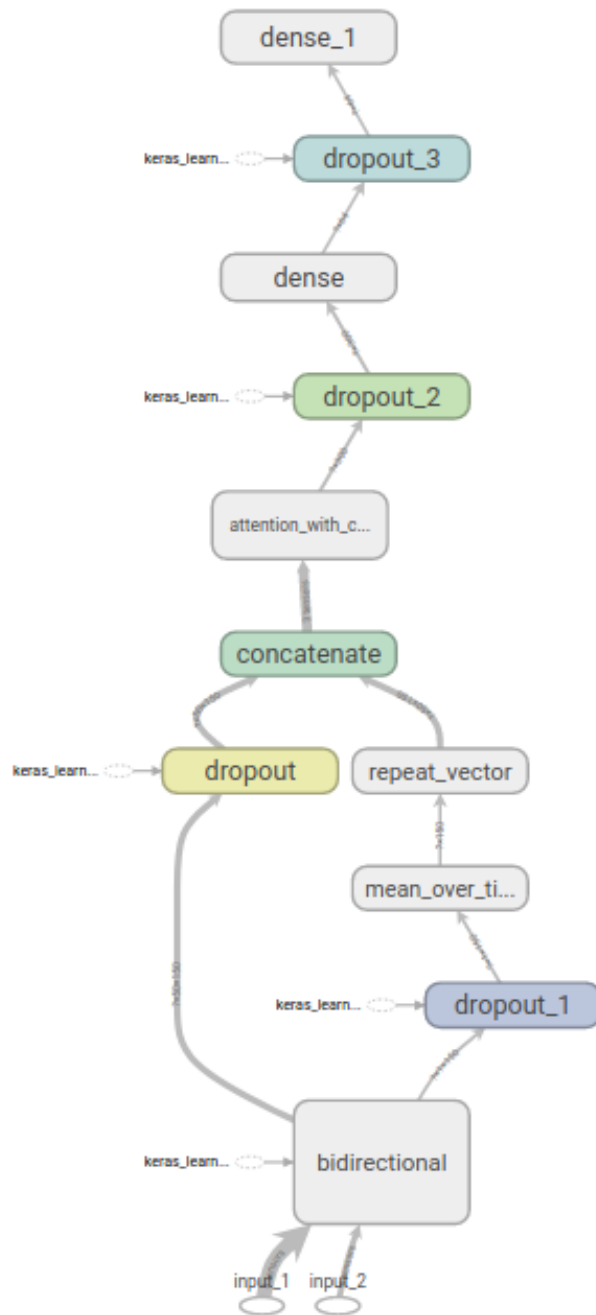


Figure 2: The model that uses AlBERTo

positive and negative class, to 3 softmax output neurons for positive, negative and mixed class. I used class weights, which help dealing with imbalanced datasets by weighting more the misclassified prediction of a class with a limited number of example. In my case the class weights used were 1.26 for negative, 1.0 for positive and 8.14 for mixed. To generate the input for the model, I extracted tuples from the dataset with the form (sentiment, topic, review).

2.1.1 Word2Vec

Initially I made some experiments using a word embedding whose dimension was 128 (<http://www.italianlp.it/resources/italian-word-embeddings/>), but this lead to a slow and limited training, that couldn't overtake the accuracy score of 0.55. So I moved to a more informative embedding whose size was 300, and although the poor quality of most of the entries (lot of words are typo or sequence of random characters), it gave good results. The Italian Word2Vec embeddings (<https://arxiv.org/pdf/2001.09332.pdf>) are obtained from a simple 2-level neural network with one-hot vector as input. It has been trained with CBOW (continuous bag of words) and negative sampling, the dataset was obtained using the information extracted from a dump of the Italian Wikipedia (dated 2019.04.01), from the main categories of Italian Google News (WORLD, NATION, BUSINESS, TECHNOLOGY, ENTERTAINMENT, SPORTS, SCIENCE, HEALTH), and has a dimension of 2.4 GB . Embeddings are given in a txt file, but after the first loading, they are formatted as dictionary and saved in a pickle file to speed up following loading. I didn't do a lot of preprocessing of the reviews, since the embedding size is big (667564 words once removed the non-ascii ones) and it could cover the 85% of the words, but looking better at the words that were indexed as `< unk >`, a lot of them were words with a capital letter. So by setting all chars lowercase, without losing much information because no proper name was present in the dataset, and just by removing all 'l', the coverage of word embeddings increased to 95%. I selected the best model by looking at the highest validation score that I obtained. I also tried to select the best model looking at the highest validation accuracy, but the results where something like 8 percentage points less. Orange is for train, blue is for validation.

2.1.2 AlBERTo

AlBERTo is the Italian BERT . It has the same structure of BERT, the difference is in the learning strategy: BERT is trained with "masked learning" and "next following sentence", while alBERTo is trained only with "masked learning". Because of this, alBERTo is suitable for sentiment analysis tasks, but not for more complex tasks like QA because it does not have idea of the flow of a dialogue. The dataset used to train alBERTo is formed by 200 million of tweets taken from TWITA, which is a big collection of tweets from 2012 since today. I used Huggingface Transformers to load the alBERTo model and the tokenizer, then I decided to generate an "alBERTed" version of the dataset, in




Figure 3: Loss for W2V on task ACP




Figure 4: Accuracy for W2V on task ACP




Figure 5: Recall for W2V on task ACP

this way I don't need to calculate the embeddings during the training because I have already computed them, and this result in a faster training at the cost of spending time and space once and for all to generate and save the embedded dataset. The alBERTo tokenizer managed to index all the words. The dimension of these embeddings is 768, which is much larger than the Word2Vec ones, and also this lead to better results and faster learning. Orange is for train, blue is for validation.




Figure 6: Loss for alBERTo on task ACP




Figure 7: Accuracy for alBERTo on task ACP




Figure 8: Recall for alBERTo on task ACP

2.2 ACD

Even if this model was not designed for topic-detection analysis, I wanted to test the potential of context-aware attention. So I proceed modifying the dataset in such a way to only know if a topic is dealt in a review, and not in which way it is dealt. The idea is that the context-aware attention should learn which words refers in a general way to a specific topic. I also had to create negative samples for the ACD train and test set, since obviously I only had positive samples. In order to do so, for each review r and for each absent topic t I added the sample ("negative", t , r) with a certain probability. I tried various probability to add a negative sample for each review, till I obtained both class weights around 1.0, hoping that a balanced dataset could give better results. The final probability I used is 0.2. The approach is the same as for the ACP task, but here the output layer has dimension 1 and so has a sigmoid activation instead of a softmax.

2.2.1 Word2Vec

The input is a list of 2 numpy array: one that contains the indexed reviews and the other one that contains the indexed topics. The output is a numpy array with values between $[0, 1]$ where 0 stands for "topic not dealt in this review" and 1 stands for "topic dealt in this review". Orange is for train, blue is for validation.




Figure 9: Loss for W2V on task ACD




Figure 10: Accuracy for W2V on task ACD




Figure 11: Recall for W2V on task ACD

2.2.2 ALBERTo

The input is a list of 2 numpy array: one that contains the embedded reviews and the other one that contains the embedded topics. The output is a numpy

array with values between $[0, 1]$ where 0 stands for "topic not dealt in this review" and 1 stands for "topic dealt in this review". Orange is for train, blue is for validation.




Figure 12: Loss for alBERTo on task ACD




Figure 13: Accuracy for alBERTo on task ACD




Figure 14: Recall for alBERTo on task ACD

3 Results

I had the official gold test set during training, but I didn't use it for model selection since that would have been "tricky". So the model selection is based on the highest validation recall. I used the official evaluation_absita script and the official gold test set for calculate the scores. The official evaluation script gives us also the macro measures, but the competition was based on micro measures.

3.1 ACD

The obvious way to predict the presence of topic in a sentence, is to try each topic for each sentence and collect the positive results only. As we can see, the

Table 2: independent results for ACP task

model	Micro-Precision	Micro-Recall	Micro-F1-score
absita best model	0.8397	0.7837	0.8108
W2V	0.3288	0.3288	0.4745
alBERTo	0.1604	0.6698	0.2588

model is not good in this task. Looking better at the results, I noticed that the model did not really understand what I was expecting, instead the output depends only by the topic and not by the review.

3.2 ACP

This results could be calculated in two ways: using the output of the ACD task as input, or use true topic presence information. In the first way, the results obviously is dependent from the result of the first task, in particular the score cannot be higher than what we obtain from the ACD task. So I only considered the independent tasks, because I want to consider the power of the model in the topic based sentiment analysis, and not in the aspect category detection. We

Table 3: INDEPENDENT results for ACP task

model	Micro-Precision	Micro-Recall	Micro-F1-score
absita best model	0.8264	0.7161	0.7673
W2V	0.9067	0.8973	0.9019
alBERTo	0.0946	0.0939	0.0942

outperform the absita best model with both embeddings by a score of more than 8 percentage points, which is a huge improvement. The alBERTo embedding improves the results obtained by the W2V one by more than 4 percentage points, and not only it has better generalization capabilities, but it learns way faster.

4 Conclusion

This model has difficulty to understand which topic is dealt in the review. This could be due to a small dataset, which contains few element for each topic class. Simpler statistical approaches like LinearSVC gives us better results, around 90% for each class, thanks to lemmatization which reduces the vocabulary and has fewer elements to analyze. Maybe some modification in both model and data could help the model to effectively learns the ACD task. Nevertheless, the model is very good at understanding the sentiment expressed toward a certain topic, and it can also understand if the sentiment is mixed, and if different topics have different sentiments in the same review. The role of the context attention is crucial in the topic based sentiment analysis, because it can understand well which words express a certain sentiment towards a certain topic. Also, alBERTo embeddings are much better than W2V one, in both test results and learning speed (alBERTo started overfitting after 20 epochs, w2v after 70 epochs) This could depend on the bigger embedding dimension (768 vs 300), the model used to obtain them (multiple transformers vs simple neural network) and the learning strategy (CBOW vs masked learning).